



University of
Nottingham

UK | CHINA | MALAYSIA

Computer Arithmetic

Dr. Heng Yu

AY2022-23, Spring Semester
COMP1047: Systems and Architecture
Week 2



Recall from your CSF

What is the decimal value of the binary number 011_2 ?

What is the binary value of the decimal number 8_{10} ?

What is the binary value of the decimal number -8_{10} ?



- **Number Formats**
- **Representing Negative Numbers**
 - **Overflow**
- **Shift and multiplication**
- **Floating Point Numbers**





After this lecture, you should be able to

- Understand and implement the format conversion between decimal and binary formats
- Representing Negative Numbers
 - Know sign/magnitude and 2's complement formats to represent a negative binary number, and sign extension
 - Understand their pros and cons.
 - Implement the related computation and conversions.
 - Identify the overflow conditions.
- Shift and multiplication
 - Implement arithmetic and logical shift. Implement binary multiplications.
- Floating Point Numbers
 - Know scientific notation, IEEE754 standard
 - Implement and convert single-precision binary values.
 - Know the special cases.





- **Number Formats**
- **Representing Negative Numbers**
 - **Overflow**
- **Shift and multiplication**
- **Floating Point Numbers**





Binary Counting

Decimal	Binary	Decimal	Binary	Decimal	Binary
0	0	6	110	12	1100
1	1	7	111	13	1101
2	10	8	1000	14	1110
3	11	9	1001	15	1111
4	100	10	1010	16	10000
5	101	11	1011	17	10001

- In binary, the result of $1 + 1$ is 0 and carries a 1 to the next digit



Binary to Decimal (Integers)

- Each binary digit corresponds to a power of **2**

Place	7 th	6 th	5 th	4 th	3 rd	2 nd	1 st	0 th
Weight	2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
	128	64	32	16	8	4	2	1

- Where the digit is 1, we add the corresponding weight
- **Example: Convert $1100\ 1010_2$ into decimal**

$$\begin{aligned} 1100\ 1010_2 &= 1 \times 128 + 1 \times 64 + 0 \times 32 + 0 \times 16 \\ &\quad + 1 \times 8 + 0 \times 4 + 1 \times 2 + 0 \times 1 \\ &= 128 + 64 + 8 + 2 = 202_{10} \end{aligned}$$



Binary to Decimal (Fractional)

- Each binary digit corresponds to a power of -2

Place	-1 st	-2 nd	-3 rd	-4 th	-5 th	-6 th	-7 th	-8 th
Weight	2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-6}	2^{-7}	2^{-8}
	0.5	0.25	0.125	0.0625	0.03125	0.015625	0.0078125	0.00390625

- Where the digit is 1, we add the corresponding weight
- **Example: Convert 0.101_2 into decimal**
- **$0.101_2 = 1 \times 2^{-1} + 1 \times 2^{-3} = 0.625_{10}$**



Binary to Decimal (Integer + Fractional)

- Convert the integer and fractional numbers separately.

Example: Convert 1011.101_2 to Decimal

$$= 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 + 1 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3}$$

$$= 8 + 0 + 2 + 1 + 0.5 + 0 + 0.125$$

$$= 11.625_{10}$$



Decimal to Binary (Integers)

- Repeatedly divide by 2, until we reach 0
- The first remainder calculated is placed at rightmost for the binary number.
- **Example: Convert 101_{10} into binary**

101	Remainder
50	1
25	0
12	1
6	0
3	0
1	1
0	1

- $101_{10} = 1100101_2$



Decimal to Binary (Fractional)

- Repeatedly multiply by 2, until we reach x.0, or infinity
- The first integer calculated is placed at leftmost after the decimal point.
- **Example: Convert 0.75_{10} into binary**

- $0.75_{10} = 0.11_2$



Decimal to Binary (Fractional)

- Example: Convert 0.3_{10} into binary

	product	integer part
0.3×2	0.6	0
0.6×2	1.2	1
0.2×2	0.4	0
0.4×2	0.8	0
0.8×2	1.6	1
0.6×2	1.2	1
0.2×2	0.4	0
0.4×2	0.8	0
0.8×2	1.6	1

...

...

...

Hence $0.3_{10} = 0.0 \ 1001 \ 1001 \ 1001 \ 1001 \dots_2$

- Digits in fraction part repeats forever, 0.3 cannot be expressed exactly by finite digits



Decimal to Binary (Integer + Fractional)

- Convert the integer and fractional numbers separately.

Convert the decimal number 95.125_{10} to binary

$$95_{10} = ?_2$$

$$0.125_{10} = ?_2$$

$$95.125_{10} = ? . ?_2$$



Decimal to Binary (Integer + Fractional)

- Convert the integer and fractional numbers separately.

Convert the decimal number 95.125_{10} to binary

$$95_{10} = 1011111_2$$

$$0.125_{10} = 0.001_2$$

$$95.125_{10} = 1011111.001_2$$



- Number Formats
- Representing Negative Numbers
 - Overflow
- Shift and multiplication
- Floating Point Numbers





Sign/Magnitude Representation

- So far, we've dealt with unsigned numbers
 - How are negative numbers represented on a computer?
- One way to represent negative number is called **sign/magnitude**.
- For an N-bit binary number:

1 sign bit, N-1 magnitude bits

Sign bit is the most significant bit (leftmost)

- Negative number: sign bit = 1
 - Positive number: sign bit = 0
-
- Rest of the bits are **numerical value** of the number



Sign/Magnitude Representation

1 sign bit, N-1 magnitude bits

Sign bit is the most significant bit (leftmost)

- Rest of the bits are **numerical value** of the number
- Hence, in a binary number with eight bits, the magnitude can range from 0000000 (0) to 1111111 (127)
- Thus numbers ranging from -127_{10} to $+127_{10}$ can be represented, once the sign bit (the eighth bit) is added

Problems

- Two representations of zero [0000, 1000]
- Difficulties in arithmetic operations:
 - $1101 (-5) + 0011 (3) = 0000$



2's Complement Representation

The most significant bit still indicates the sign

Same as unsigned binary, but the value of the most significant bit is -2^{N-1}

Example: What is the value of 10110

- Unsigned: $10110 = 16 + 4 + 2 = 22$
- Twos complement $10110 = -16 + 4 + 2 = -10$

Single representation of zero (0000)

Arithmetic works fine

- $1011 (-5) + 0011 (3) = 1110 (-8 + 4 + 2) = -2$



2's Complement Representation

How to convert a 2's complement binary to its decimal value?

- Assume in an 8-bit two's-complement numeral system

8-Bit Two's Complement ($-128 \leq x < 127$)

Bit	MSB							LSB
	7 th	6 th	5 th	4 th	3 rd	2 nd	1 st	0 th
Weight	-2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0

- What does 0000 0001 represent? **1**
- What does 1111 1111 represent? **-1**
- What does 0101 1011 represent? **91**



2's Complement Representation

How to convert a positive 2's complement binary to its negative value?

E.g., what is the 2's complement of the binary number 011010

First method

Complement each digit to get the first complement

100101

Add one to the first complement.

100110

Second Method

Start from the least significant bit, locate the first digit with value 1

011010

Complement each digit after the first digit with value 1.

100110



2's Complement Representation

- Exercise: Represent -27_{10} into 2's complement binary format



2's Complement Representation

- The two's-complement numeral system is **the most common method** of representing signed integers on computers
 - Advantages: The **fundamental arithmetic operations** of addition, subtraction, and multiplication are identical to those for unsigned binary numbers
 - This property makes the system both simpler to implement and capable of easily handling higher precision arithmetic
 - Also, **zero has only a single representation** 0, eliminating the subtleties associated with negative zero
- All modern processors primarily use two's complement
 - Includes all the MIPS's integer arithmetic instructions
 - Different instruction variants for signed and unsigned



Sign Extension

- Sign extension is the operation of increasing the number of bits of a binary number while **preserving the number's sign (positive/negative) and value**
 - This is done by appending digits to the most significant side of the number, following a procedure dependent on the particular signed number representation used
- In MIPS, **all arithmetic immediate values are sign-extended**
- e.g., If six bits are used to represent the number "00 1010" (decimal +10) and the sign extension operation increases the word length to 16 bits, then the new representation is simply "0000 0000 0000 1010" – padding the left side with 0s



Sign Extension

- Question: If ten bits are used to represent the value "11 1111 0001" (decimal -15) in two's complement, and the sign extension operation increases the word length to 16 bits, what is the new representation?



Sign Extension

- **Solution:** If ten bits are used to represent the value "11 1111 0001" (decimal -15) in two's complement, and the sign extension operation increases the word length to 16 bits, the new representation is "1111 1111 1111 0001" – padding the left side with 1s

Long Addition in Decimal

	6	2	9	5	1	4	1	3	
+	2	8	1	8	2	8	1	7	
	0	1	1	1	0	1	0	1	Carry
	9	1	1	3	4	2	3	0	

0111 0110 + 1101 0101

Long Addition in Binary

	0	1	1	1	0	1	1	0	= 118
+	1	1	0	1	0	1	0	1	= 213
	1	1	1	1	0	1	0	0	Carry
<hr style="border: 0.5px solid black;"/>									
	1	0	1	0	0	1	0	1	= 331




- One issue in computer arithmetic is dealing with **finite amounts of storage**, such as 32-bit MIPS registers
- Overflow occurs when the result of an operation is too large to be stored
- For an unsigned number, overflow happens when the last carry (1) cannot be accommodated. **E.g. 1111+0001**



- For a signed number, overflow occurs when
 - Adding two positives yields a negative. **E.g. (0111 + 0001)**
 - Or, adding two negatives gives a positive. **E.g. (1000 + 1000)**
 - Or, subtract a negative from a positive gives a negative
 - Or, subtract a positive from a negative gives a positive
- One way to detect overflow is to check **whether the sign bit is consistent** with the sign of the inputs when the two inputs are of the same sign – if you added two positive numbers and got a negative number, something is wrong, and vice versa

Ariana 5

- In 1996, the European Space Agency's Ariane5 rocket was launched for the first time... and it exploded 40 seconds after liftoff
- It turns out that the Ariane5 **used software designed for the older Ariane4**
 - The Ariane4 stored its horizontal velocity as a 16-bit signed integer
 - But the Ariane5 reaches a much higher velocity, which **caused an overflow** in the 16-bit quantity
- The overflow error was never caught, so incorrect instructions were sent to the rocket boosters and main engine
- For a modern CPU like MIPS, it detects overflow with an exception (an internal interrupt signal to the CPU)
 - Control jumps to predefined address for exception
 - Interrupted address is saved for resumption





- Number Formats
- Representing Negative Numbers
 - Overflow
- **Shift and multiplication**
- Floating Point Numbers





Shift Operations

- Shift operations shift a word **a number of places to the left or right**
- Bits which are shifted out, just disappear
- E.g. on 4 bits: **1011** shifted left 1 bit, results in **0110**.
- **Logical shift**: in **right** shift, **0s** are filled in empty positions.
 - E.g., 1011 shift right 1 bit => 0101
- **Arithmetic shift**: in **right** shift, **sign bits** are filled in empty positions.
 - E.g., 1011 shift right 1 bit => 1101



Shift Operations

- For unsigned and 2's complement numbers, **logical/arithmetic shift left** by 1 is **multiplication** by 2 if there is no overflow
 - e.g. 0011 (3) shifted left 1 bit results in 0110 (6)
 - e.g. 1100 (-4) shifted left 1 bit results in 1000 (-8)
- For unsigned numbers, **logical shift right** by 1 is **division** by 2, ignoring remainder.
 - e.g. 0101 (5) shifted right 1 bit results in 0010 (2)
 - For 2's complement numbers, logical shift right by 1 does not correspond to dividing 2.
 - e.g. 1100 (-4) shifted right results in 0110 (6)
- For 2's complement numbers, **arithmetic shift right** by 1 performs **division** by 2.
 - e.g. 1100 (-4) shifted right arithmetical results in 1110 (-2)



Multiplication

- Binary multiplication similar to decimal multiplication – multiplying with each bit and add the (shifted) results together.
- If we multiply an m -bit with an n -bit number, we need $m+n$ bits to store the result
 - Multiplying 4-digit numbers needs 8 digits
 - e.g. $1101_2 \times 1011_2$

Long Multiplication in Binary

					1	1	0	1	$a = 13_{10}$
					1	0	1	1	$b = 11_{10}$
					<hr/>				
					1	1	0	1	
				1	1	0	1		
			0	0	0	0			
		0	0	0	0				
	+		1	1	0	1			Partial Sums
	<hr/>								
	=	1	0	0	0	1	1	1	$c = 143_{10}$



- Number Formats
- Representing Negative Numbers
 - Overflow
- Shift and multiplication
- **Floating Point Numbers**





- Scientific notation
 - A real number notation that renders with a **single** digit to the left of the **decimal** point
- General form: $\pm m \times b^e$
 - ***m*** is called the significand
 - ***b*** is the base number, can be any positive integer
 - Exponent ***e*** is an integer
- E.g.
 - 2.73×10^7
 - 0.273×10^{-6}




Normalized Scientific Notation

- Normalized scientific notation
 - The exponent e is chosen so that the absolute value of m remains at least one but less than ten ($1 \leq |m| < 10$)
 - A number in proper scientific notation that has no leading zeros
- E.g.
 - 2.73×10^7
 - But 0.273×10^{-6} is not normalized
- Normalized scientific notation for binary
 - A number in floating-point notation that has no leading 0s
 - E.g. $11.101_2 = 11.101_2 \times 2^0$
 $= 1.1101_2 \times 2^1$ (normalized form)
 - Fraction: the value placed in the fraction field (1101 in this case)



Normalized Scientific Notation



- When we normalize a non-zero binary number, we'll **always have a 1** to the left of the binary point
- To represent a real (floating point) number in a **fixed number of digits**, we need to decide how to allocate some fixed number of bits to both the fraction *xxxx* and the exponent *yyyy*
 - This is a tradeoff between precision and range



Floating Point Standard

- The IEEE Standard for Floating-Point Arithmetic (**IEEE 754**) is a technical standard for floating-point computation established in 1985
- It gives exact layout of bits, and defines basic arithmetic
- Two representations
 - Single precision (32-bit)
 - Double precision (64-bit)
- Now almost universally adopted




IEEE 754 Floating Point Standard

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{BiasedExponent} - \text{Bias})}$$


- A type of sign and magnitude representation
- S: sign bit of the fraction (0 for +, 1 for -)



- Single-precision (**32-bit**)



- Double-precision (**64-bit**)





IEEE 754 Floating Point Standard

S	Biased Exponent	Fraction
---	--------------------	----------

- Biased Exponent
 - A fixed value, called bias, is subtracted from the field to get the true exponent value
 - Typically, $\text{bias} = 2^{k-1} - 1$, where k is the number of bits in the exponent field
 - E.g. the biased representation $0000\ 0101_2$ with a bias 127, its actual exponent value is: $5 - 127 = -122$ ($00000101_2 = 5_{10}$)
- Biased Exponent = Actual Exponent + Bias
- IEEE 754 Single precision, bias=127, double precision, bias=1023



A Single Precision Example

- Represent the following binary floating point in the Single-Precision Floating Point Format
 - -10.101_2
- Step 1: normalize number $(-1) \times 1.0101 \times 2^1$
- Step 2:
 - Sign: 1, fraction $F=0101$
 - Biased Exponent $E = \text{Actual Exponent} + \text{Bias}$
 $= 1 + 127 = 128_{10} = 1000\ 0000_2$

8 bits		23 bits			
1	$E=1000\ 0000$	$F=0101$	0000	0000	0000 0000 000



Try By Yourself

- Represent the following binary floating point in the Single-Precision Floating Point Format
 - 1011.1101_2
- Step 1: normalize number $(1) \times 1.0111101 \times 2^3$
- Step 2:
 - Sign: 0, fraction $F=0111101$
 - Biased Exponent $E = \text{Actual Exponent} + \text{Bias}$
 $= 3 + 127 = 130_{10} = 1000\ 0010_2$

8 bits		23 bits					
0	$E=1000\ 0010$	$F=0111$	1010	0000	0000	0000	000



Try By Yourself

- Represent the following binary floating point in the Single-Precision Floating Point Format
 - 1011.1101_2



Special Values (Single Precision)

- Smallest positive normalized number: $1.0_2 \times 2^{-126}$
- Least negative normalized number: $-1.0_2 \times 2^{-126}$
- Zero: $\pm 1.0_2 \times 2^{-127}$
- Largest positive normalized number:
 $1.1111\ 1111\ 1111\ 1111\ 1111\ 111_2 \times 2^{127}$
- Most negative normalized number:
 $-1.1111\ 1111\ 1111\ 1111\ 1111\ 111_2 \times 2^{127}$
- $\pm\infty$: $\pm 1.0 \times 2^{128}$
- NaN: $1.f \times 2^{128}$ where $f \neq 0$



NaN (Not-a-Number)

- A numeric data type value representing an undefined or unrepresentable value, especially in floating-point calculations
- A few cases where we get NaN:
 - $0.0 / 0.0$
 - $\pm\infty / \pm\infty$,
 - $0 \times \pm\infty$,
 - $-\infty + \infty$,
 - $\text{sqrt}(-1.0)$,
 - $\text{log}(-1.0)$

```
void nanFun()
{
    printf("0.0/0.0: %f\n", 0.0/0.0);
    printf("inf/inf: %f\n", (1.0/0.0)/(1.0/0.0));
    printf("0.0*inf: %f\n", 0.0*(1.0/0.0));
    printf("-inf + inf: %f\n", (-1.0/0.0) + (1.0/0.0));
    printf("sqrt(-1.0): %f\n", sqrt(-1.0));
    printf("log(-1.0): %f\n", log(-1.0));

    float n = log(-1.0);
    printBits(sizeof(n), &n);
}
```



- Convert the following decimal number to single-precision floating point representation in binary format

$$x = -13.825$$

- Hint: first convert to the binary correspondence, then change it into single-precision floating point representation



Putting It All Together

$$x = -13.825$$

Integer part: $13_{10} = 1101_2$, fraction part:

0.825	integer
1 .65	1
1 .3	1
0.6	0
1 .2	1
0.4	0
0.8	0
1 .6	1
1 .2	1
<i>repeats ...</i>	

therefore, $x = 1101.1101001_2 = 1.1011101001 \times 2^3_2$.

Exponent: $3 + 127 = 130 = 1000\ 0010$.

8 bits		23 bits					
1	E=1000 0010	F=101	1101	0011	0011	0011	0011



Equality Conditions

- In most computers, $0.3+0.2 \neq 0.5$
 - Neither 0.3 nor 0.2 has exact binary representation
- Never directly test floating point numbers for equality (i.e. $0.3+0.2 == 0.5$)
 - Using the following statement instead
$$|0.3+0.2-0.5| < \epsilon$$
where ϵ is a very small number
- You don't have to worry about it as a user. MIPS instructions already implement the details.



- Conversion between binary and decimal numbers
- Representing negative numbers
 - Sign/magnitude
 - 2's complement
 - Sign extension
- Binary additions and overflow, multiplication
- Shift operations
 - Logical shift
 - Arithmetic shift
- Floating point number representation.
 - IEEE 754 standard for single precision.
 - Special values.



University of
Nottingham

UK | CHINA | MALAYSIA

Stay Tuned.