

# La Trobe University

## Department of Computer Science and Computer Engineering

### CSE3BDC Assignment 2023

#### Objectives

1. Gain in depth experience playing around with big data tools (Hive, SparkRDDs, and Spark SQL).
2. Solve challenging big data processing tasks by finding highly efficient solutions.
3. Experience processing three different types of real data
  - a. Standard multi-attribute data (Bank data)
  - b. Time series data (Twitter feed data)
  - c. Bag of words data.
4. Practice using programming APIs to find the best API calls to solve your problem. Here are the API descriptions for Hive, Spark (especially spark look under RDD. There are a lot of really useful API calls).
  - a) [Hive] <https://cwiki.apache.org/confluence/display/Hive/LanguageManual>
  - b) [Spark] <http://spark.apache.org/docs/latest/api/scala/index.html#package>
  - c) [Spark SQL] <https://spark.apache.org/docs/latest/sql-programming-guide.html>  
<https://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.sql.Database>  
<https://spark.apache.org/docs/latest/api/sql/index.html>

- If you are not sure what a spark API call does, try to write a small example and try it in the spark shell

**This assignment is due 23:59 p.m. on Friday 19<sup>th</sup> of May, 2023.**

Penalties are applied to late assignments (accepted up to 5 business days after the due date only). Five percent is deducted per business day late. A mark of zero will be assigned to assignments submitted more than 5 days late.

This is an **individual** assignment. You are not permitted to work as a part of a group when writing this assignment.

#### Submission checklist

- Ensure that all of your solutions read their input from the full data files (not the small example versions)
- Check that all of your solutions run without crashing in the docker containers that you used in the labs.
- Delete all output files
- Archive up everything into a single zip file and submit your assignment via LMS

## Copying, Plagiarism

Plagiarism is the submission of somebody else's work in a manner that gives the impression that the work is your own. For individual assignments, plagiarism includes the case where two or more students work collaboratively on the assignment. The Department of Computer Science and Computer Engineering treats plagiarism very seriously. When it is detected, penalties are strictly imposed.

**If you are working on your assignment on the lab computers. Make sure you delete the virtual machine and empty the recycle bin before you leave. Otherwise, other students may be able to see your solutions.**

## ChatGPT and similar AI tools

A key purpose of this assessment task is to test your own ability to complete the assigned tasks. Therefore, the use of ChatGPT, AI tools or chatbots with similar functionality is prohibited for this assessment task. Students who are found to be in breach of this rule will be subject to normal academic misconduct measures. Additionally, students may be engaged to provide an oral validation of their understanding of their submitted work (e.g. coding).

## Expected quality of solutions

- a) In general, writing more efficient code (less reading/writing from/into HDFS and less data shuffles) will be rewarded with more marks.
- b) This entire assignment can be done using the docker containers supplied in the labs and the supplied data sets **without running out of memory**. It is time to show your skills!
- c) I am not too fussed about the layout of the output. As long as it looks similar to the example outputs for each task. That will be good enough. The idea is not to spend too much time massaging the output to be the right format but instead to spend the time to solve problems.
- d) For Hive queries. We prefer answers that use less tables.

The questions in the assignment will be labelled using the following:

- [Hive]
  - Means this question needs to be done using Hive
- [Spark RDD]
  - Means this question needs to be done using Spark RDDs, you are **not allowed** to use any Spark SQL features like dataframe or datasets.
- [Spark SQL]
  - Means this question needs to be done using Spark SQL and therefore you are not allowed to use RDDs. In addition, you need to do these questions using the spark dataframe or dataset API, do not use SQL syntax.

## Assignment structure:

- A script which puts all of the data files into HDFS automatically is provided for you. Whenever you start the docker container again you will need to run the following script to upload the data to HDFS again, since HDFS state is not maintained across docker runs:

```
$ bash put_data_in_hdfs.sh
```

The script will output the names of all of the data files it copies into HDFS. If you do not run this script, solutions to the Spark questions **will not work** since they load data from HDFS.

- For each Hive question a skeleton .hql file is provided for you to write your solution in. You can run these just like you did in labs:

```
$ hive -f Task_XX.hql
```

- For each Spark question, a skeleton project is provided for you. Write your solution in the .scala file in the src directory. Build and run your Spark code using the provided scripts:

```
$ bash build_and_run.sh
```

## Tips:

1. Look at the data files **before** you begin each task. Try to understand what you are dealing with!
2. For each subtask we provide small example input and the corresponding output in the assignment specifications below. These small versions of the files are also supplied with the assignment (they have “-small” in the name). It’s a good idea to get your solution working on the small inputs first before moving on to the full files.
3. In addition to testing the correctness of your code using the very small example input. You should also use the large input files that we provide to test the scalability of your solutions.
4. It can take some time to build and run Spark applications from .scala files. So for the Spark questions it’s best to experiment using *spark-shell* first to figure out a working solution, and then put your code into the .scala files afterwards.

## Task 1: Analysing Bank Data [38 marks total]

We will be doing some analytics on real data from a Portuguese banking institution<sup>1</sup>. The data is stored in a semicolon (“;”) delimited format.

The data is supplied with the assignment at the following locations:

Small version	Full version
Task_1/Data/bank-small.csv	Task_1/Data/bank.csv

The data has the following attributes

Attribute index	Attribute name	Description
0	age	numeric
1	job	type of job (categorical: admin., unknown, unemployed, management, housemaid, entrepreneur, student, blue-collar, self-employed, retired, technician, services)
2	marital	marital status (categorical: married, divorced, single; note: divorced means divorced or widowed)
3	education	(categorical: unknown, secondary, primary, tertiary)
4	default	has credit in default? (binary: yes, no)
5	balance	average yearly balance, in euros (numeric)
6	housing	has housing loan? (binary: yes, no)
7	loan	has personal loan? (binary: yes, no)
8	contact	contact communication type (categorical: unknown, telephone, cellular)
9	day	last contact day of the month (numeric)
10	month	last contact month of year (categorical: jan, feb, mar, ..., nov, dec)
11	duration	last contact duration, in seconds (numeric)
12	campaign	number of contacts performed during this campaign and for this client (numeric, includes last contact)
13	pdays	number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)
14	previous	number of contacts performed before this campaign and for this client (numeric)
15	poutcome	outcome of the previous marketing campaign (categorical: unknown, other, failure, success)
16	termdeposit	has the client subscribed a term deposit? (binary: yes, no)

1 : Banking data source: <http://archive.ics.uci.edu/ml/datasets/Bank+Marketing>

Here is a small example of the bank data that we will use to illustrate the subtasks below (we only list a subset of the attributes in this example, see the above table for the description of the attributes):

<u>job</u>	<u>marital</u>	<u>education</u>	<u>balance</u>	<u>loan</u>
management	Married	tertiary	2143	Yes
technician	Divorced	secondary	29	Yes
entrepreneur	Single	secondary	2	No
blue-collar	Married	unknown	1506	No
services	Divorced	secondary	829	Yes
technician	Married	tertiary	929	Yes
Management	Divorced	tertiary	22	No
technician	Married	primary	10	No

Please note we specify whether you should use [Hive] or [Spark RDD] for each subtask at the beginning of each subtask.

- a) [Hive] Report the number of clients for each marital status who have a balance above 500 and has a loan. Write the results to Task\_1a-out. For the above small example data set you would report the following (output order is not important for this question):

```
divorced,1
married,2
```

[8 marks]

- b) [Hive] Report the average yearly balance for all people in each job category in descending order of average yearly balance. Write the results to Task\_1b-out. For the small example data set you would report the following:

```
blue-collar,1506.0
management,1082.5
services,829.0
technician,322.6666666666667
entrepreneur,2.0
```

[8 marks]

- c) [Spark RDD] Group balance into the following three categories:
- Low: -infinity to 500
  - Medium: 501 to 1500 =>
  - High: 1501 to +infinity

Report the number of people in each of the above categories. Write the results to "Task\_1c-out" in text file format. For the small example data set you should get the following results (output order is not important in this question):

High,2  
Medium,2  
Low,4

[10 marks]

- d) [Spark RDD] Output the following details for each person whose job category has an average balance above 500: *education, balance, job, marital, loan*. Make sure the output is in decreasing order of individual balance. Write the results to Task\_1d-out in text file format (output to a single file). For the small example data set you would report the following:

tertiary, 2143.0, management, married, yes  
unknown, 1506.0, blue-collar, married, no  
secondary, 829.0, services, divorced, yes  
tertiary, 22.0, management, divorced, no

[12 marks]

## Task 2: Analysing Twitter Time Series Data [32 marks]

In this task we will be doing some analytics on real Twitter data<sup>2</sup>. The data is stored in a tab (“t”) delimited format.

The data is supplied with the assignment at the following locations:

Small version	Full version
Task_2/Data/twitter-small.tsv	Task_2/Data/twitter.tsv

The data has the following attributes

Attribute index	Attribute name	Description
0	tokenType	In our data set all rows have Token type of hashtag. So this attribute is useless for this assignment.
1	month	The year and month specified like the following: YYYYMM. So 4 digits for year followed by 2 digits for month. So like the following 200905, meaning the year 2009 and month of May
2	count	An integer representing the number tweets of this hash tag for the given year and month
3	hashtagName	The #tag name, e.g. babylove, mydate, etc.

Here is a small example of the Twitter data that we will use to illustrate the subtasks below:

Token type	Month	count	Hash Tag Name
hashtag	200910	2	babylove
hashtag	200911	2	babylove
hashtag	200912	90	babylove
hashtag	200812	100	mycoolwife
hashtag	200901	201	mycoolwife
hashtag	200910	1	mycoolwife
hashtag	200912	500	mycoolwife
hashtag	200905	23	abc
hashtag	200907	1000	abc

2 : Twitter data source: <http://www.infochimps.com/datasets/twitter-census-conversation-metrics-one-year-of-urls-hashtags-sm>

- a) [Spark RDD] Find the single row that has the highest count and for that row report the month, count and hashtag name. Print the result to the terminal output using `println`. So, for the above small example data set the result would be:

month: 200907, count: 1000, tagName: abc
--

[6 marks]

- b) [Do twice, once using Hive and once using Spark RDD] Find the hash tag name that was tweeted the most in the entire data set across all months. Report the total number of tweets for that hash tag name. You can either print the result to the terminal or output the result to a text file. So, for the above small example data set the output would be:

abc 1023

[12 marks total: 6 marks for Hive and 6 marks for Spark RDD]

- c) [Spark RDD] Given two months x and y, where  $y > x$ , find the hashtag name that has increased the number of tweets the most from month x to month y. Ignore the tweets in the months between x and y, so just compare the number of tweets at month x and at month y. Report the hashtag name, the number of tweets in months x and y. Ignore any hashtag names that had no tweets in either month x or y. You can assume that the combination of hashtag and month is unique. Therefore, the same hashtag and month combination cannot occur more than once. Print the result to the terminal output using `println`. For the above small example data set:

<b>Input</b>	x = 200910, y = 200912
<b>Output</b>	tagName: mycoolwife, countX: 1, countY: 500

For this subtask you can specify the months x and y as arguments to the script. This is required to test on the full-sized data. For example:

```
$ bash build_and_run.sh 200901 200902
```

[14 marks]



## Task 3: Indexing Bag of Words data [30 marks]

In this task you are asked to create a partitioned index of words to documents that contain the words. Using this index you can search for all the documents that contain a particular word efficiently.

The data is supplied with the assignment at the following locations<sup>3</sup>:

Small version	Full version
Task_3/Data/docword-small.txt	Task_3/Data/docword.txt
Task_3/Data/vocab-small.txt	Task_3/Data/vocab.txt

The first file is called *docword.txt*, which contains the contents of all the documents stored in the following format:

Attribute index	Attribute name	Description
0	docId	The ID of the document that contains the word
1	vocabId	Instead of storing the word itself, we store an ID from the vocabulary file.
2	count	An integer representing the number of times this word occurred in this document.

The second file called *vocab.txt* contains each word in the vocabulary, which is indexed by vocabIndex from the *docword.txt* file.

Here is a small example content of the *docword.txt* file.

docId	vocabId	count
3	3	600
2	3	702
1	2	120
2	5	200
2	2	500
3	1	100
3	5	2000
3	4	122
1	3	1200
1	1	1000

Here is an example of the *vocab.txt* file

vocabId	word
1	plane
2	car
3	motorbike
4	truck
5	boat

Complete the following subtasks using Spark:

- a) [spark SQL] Calculate the total count of each word across all documents. List the words in ascending alphabetical order. Write the results to “Task\_3a-out” in CSV format (multiple output parts are allowed). So for the above small example input the output would be the following (outputs with multiple parts will be considered in order of the part number):

```
boat,2200
car,620
motorbike,2502
plane,1100
truck,122
```

Note: spark SQL will give the output in multiple files. You should ensure that the data is sorted globally across all the files (parts). So, all words in part 0, will be alphabetically before the words in part 1.

[10 marks]

- b) [spark SQL] Find the most frequently occurring word for each document and then output the following information: docId, word, count. Sort in decreasing order according to count. Write the results to Task\_3b-out in CSV format (multiple output parts are allowed).

So, for the above small example input, the output would be the following:

```
3, boat, 2000
1, motorbike, 1200
2, motorbike,702
```

Note: spark SQL will give the output in multiple files. You should ensure that the data is sorted globally across all the files (parts).

[20 marks]

## Bonus Marks:

1. Using spark perform the following task using the data set of task 2.

[Spark RDD or Spark SQL] Find the hash tag name that has increased the number of tweets the most from among any two consecutive months of any hash tag name. Consecutive month means for example, 200801 to 200802, or 200902 to 200903, etc. Report the hash tag name, the 1<sup>st</sup> month count, and the 2<sup>nd</sup> month count using `println`.

For the small example data set of task 2 the output would be:

Hash tag name: mycoolwife  
count of month 200812: 100  
count of month 200901: 201

[10 marks]

## Total Marks:

Please note that the total mark for this assignment is capped at 100. If your marks add to more than 100 then your final mark will be 100.

### **Return of Assignments**

Departmental Policy requires that assignments be returned within three weeks of the submission date. We will endeavour to have your assignment returned before the BDC exam. The time and place of the return will be posted on LMS.