MIS772

Predictive Analytics

Workshop: Intro to Classification

Decision trees and k-NN with holdout validation

DEAKIN BUSINESS SCHOOL

AACSB ACCREDITED

EFMD EQUIS ACCREDITED

Workshop Plan

Objectives:

The task is to create a predictive model classifying all Danish AirBnB rental properties into "cheap" (price/night < $100) and "expensive". Visualize properties located in the Danish region of Northern Jutland by generating a new attribute.

Data Set:

AirBNB-DK.csv (From Unit Website)

Acknowledgements: *http://tomslee.net*

Method:

Attend the workshop, follow the tutor's demo and instructions, take notes. Note that the class and online seminar will be recorded and their videos linked to the CloudDeakin topic for later access and study.

1. **Prepare data for modelling**
   (a) Load data from file
   (b) Filter out data with missing values
   (c) Generate new attributes, using function expressions
   (d) Discretize an attribute
   (e) Select predictors and a label attribute
   (f) Set role (label)
2. **Train and validate decision tree model**
   (a) Split data into training and validation partition
   (b) Add a ***k-NN*** or ***Decision Tree*** model and train it
   (c) Apply the model to validation data
   (d) Measure and interpret the model and performance
   (e) Save the RM process (use versions)
3. **Discuss, extend and reflect**

DEAKIN BUSINESS SCHOOL

Visualise geospatial data

Watch a lecture demo. Follow the tutor. Initially create an RM process responsible for data pre-processing, similar to the one below. It will load the data from a .csv file, filter out data with missing attributes, generate new attributes (including identifying North Jutland properties). Visualize properties located in North Jutland using a point map.

How can this be done?

What results does it produce?

How does it work?

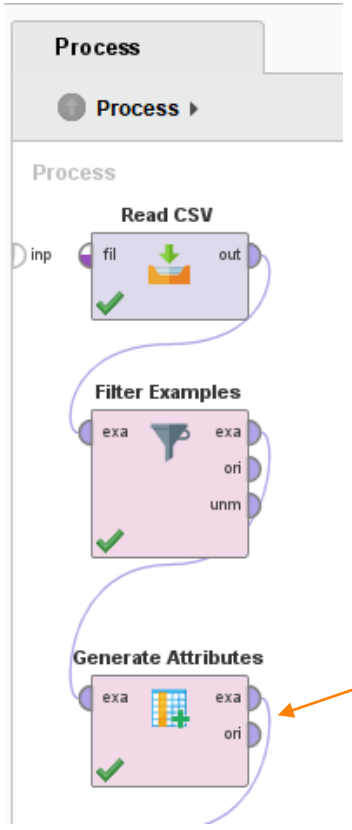

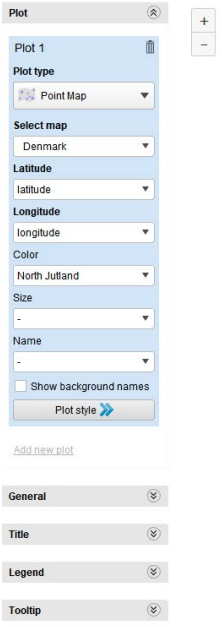

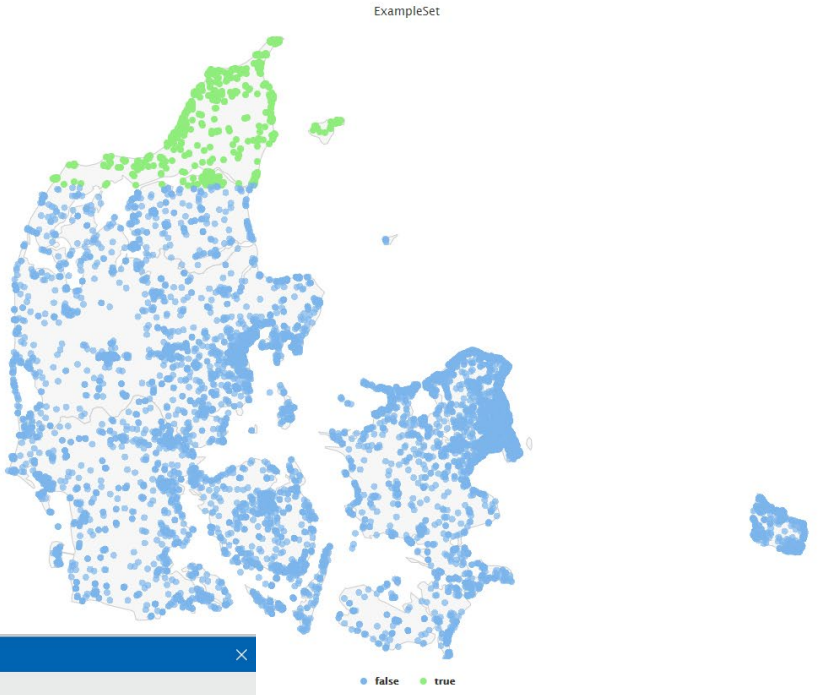

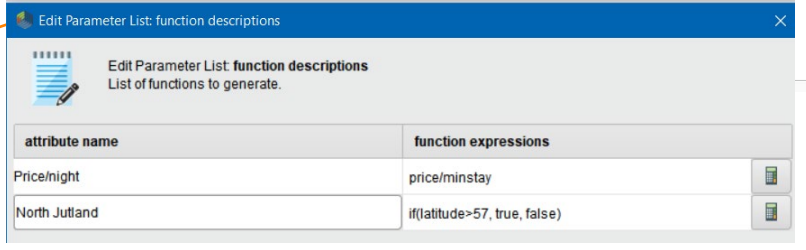

DEAKIN BUSINESS SCHOOL

Discretize attribute

Discretize the price/night as cheap (<$100) or expensive.

How can this be done?

What results does it produce?

How does it work?

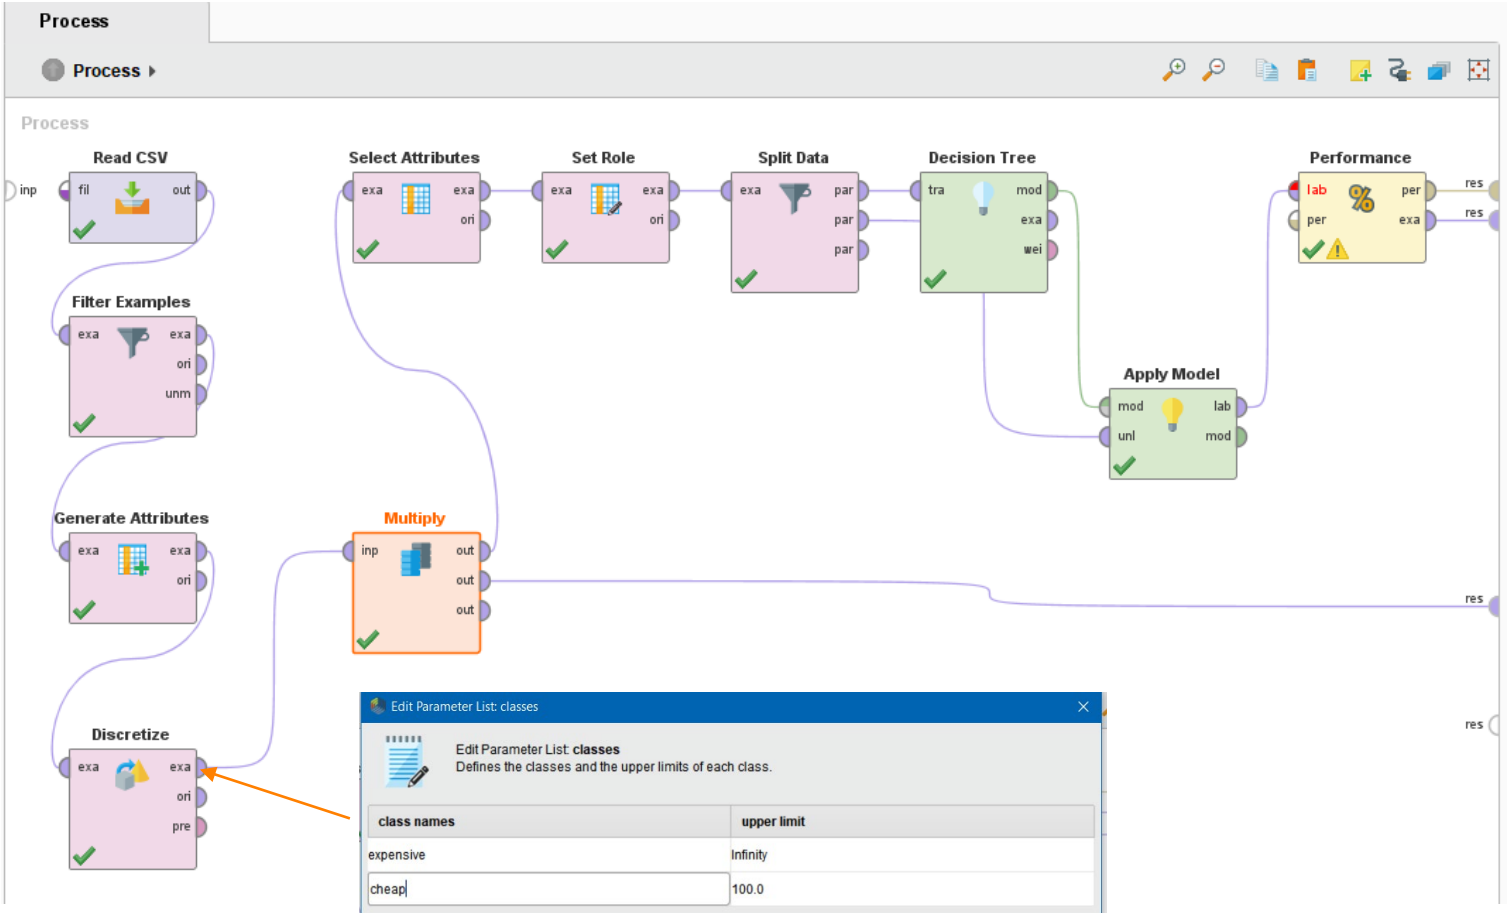

DEAKIN BUSINESS SCHOOL

Select predictor attributes and set role

Select predictor attributes and set the role of the price/night category as a label. Run and check the results.

How can this be done?

What results does it produce?

How does it work?

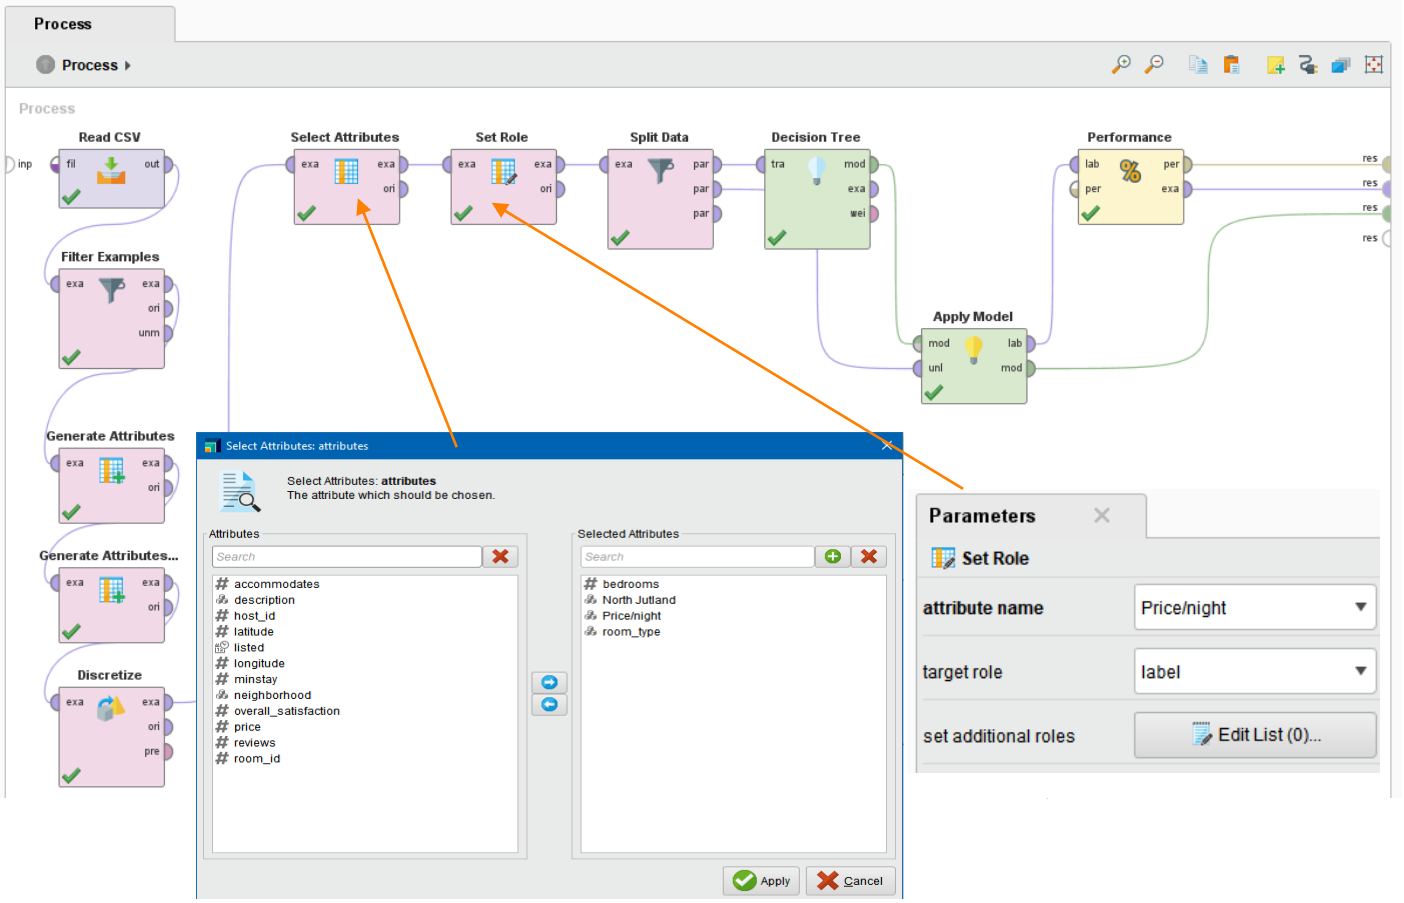

DEAKIN BUSINESS SCHOOL

The End Result: The RM Process

Next we will develop model training and validation. You will add and test one operator at a time. First split data into two parts 70-30%, next add a k-NN or a decision tree, then apply the resulting model to validation data, and finally analyse the predictions using binomial performance (accuracy and kappa). Run and check the results. Experiment with the DT parameters. Save your work.

We will extend the previous process.

What results does it produce?

How does it work?

Make sure: You do a little, test a little, think a little and save a lot!

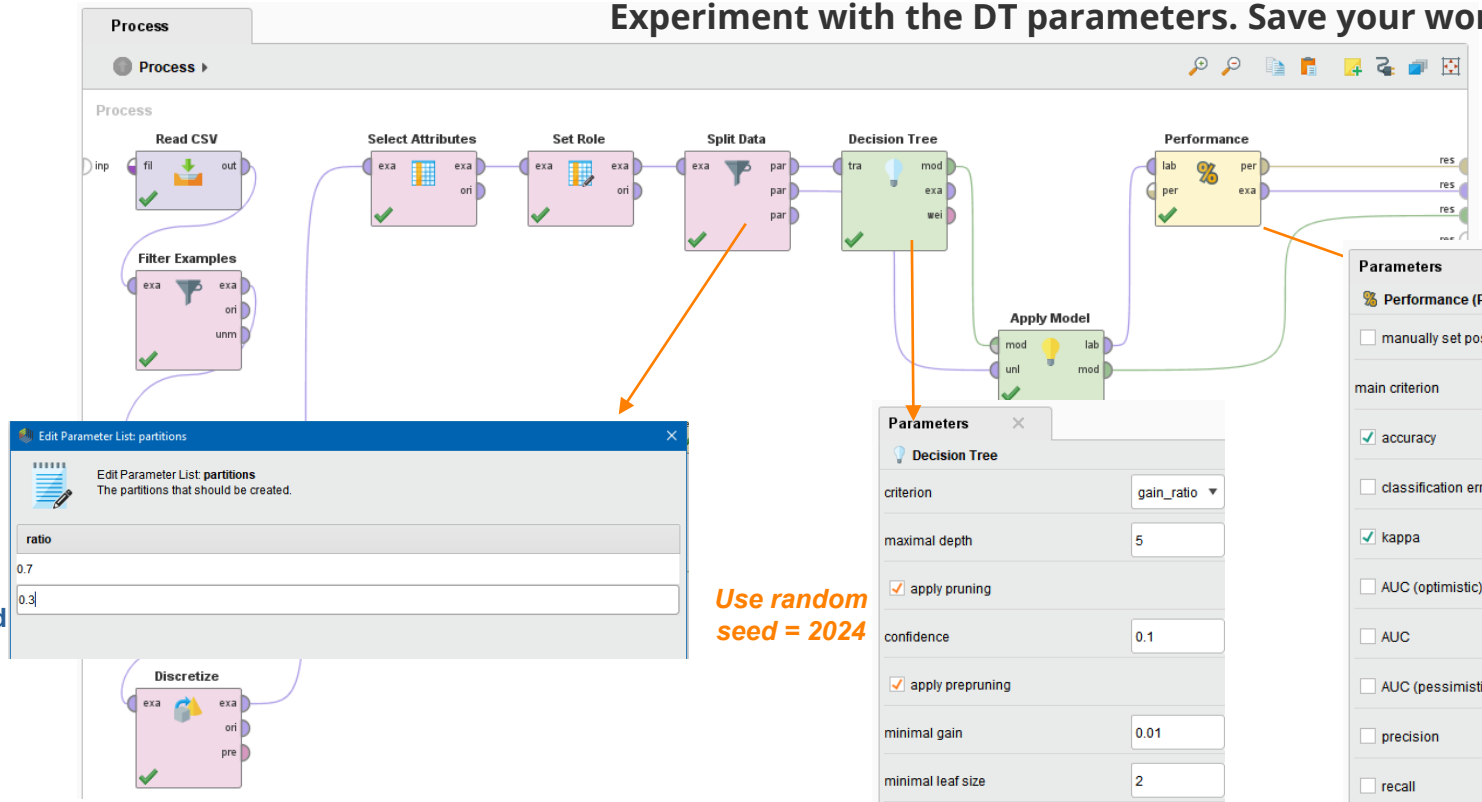

Use random seed = 2024

What changes are needed for this holdout validation to work with other classifiers e.g., Logistic regression?

DEAKIN BUSINESS SCHOOL

The End Result: Tables & Charts

Explore classification results. View the results table – inspect label vs prediction, how should we read the confidence columns. Then, look at the tree model. Analyse the model performance, is it good or bad? Can it be improved? How? How can the skills gained transfer to your assignment?

What are confidence columns and how do they relate to prediction?

Row No.	Price/night	prediction(P...	confidence(...	confidence(...	room_type	bedrooms	North Jutland
1	cheap	cheap	0.900	0.100	Private room	1	false
2	cheap	cheap	0.900	0.100	Private room	1	false
3	cheap	cheap	0.900	0.100	Private room	1	false
4	cheap	cheap	0.900	0.100	Private room	1	false
5	cheap	cheap	0.900	0.100	Private room	1	false
6	expensive	cheap	0.669	0.331	Entire home/...	0	false
7	expensive	cheap	0.669	0.331	Entire home/...	1	false
8	expensive	expensive	0.190	0.810	Entire home/...	6	false
9	cheap	cheap	0.669	0.331	Entire home/...	1	false
10	cheap	cheap	0.900	0.100	Private room	1	false
11	cheap	cheap	0.900	0.100	Private room	1	false
12	cheap	cheap	0.900	0.100	Private room	1	false
13	cheap	cheap	0.900	0.100	Private room	1	false
14	cheap	cheap	0.900	0.100	Private room	1	false
15	cheap	cheap	0.900	0.100	Private room	1	false
16	cheap	cheap	0.900	0.100	Private room	1	false
17	cheap	cheap	0.900	0.100	Private room	1	false
18	expensive	expensive	0.381	0.619	Entire home/...	3	false
19	cheap	cheap	0.900	0.100	Private room	1	false

View the decision tree visualisation. Experiment with the model parameters

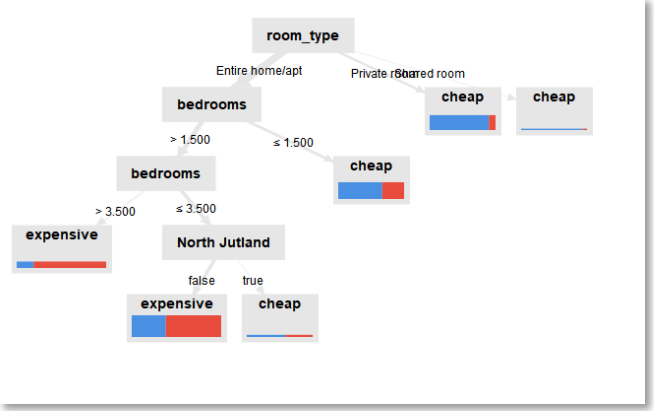

What does it mean?

PerformanceVector

```
PerformanceVector:
accuracy: 72.69%
ConfusionMatrix:
True:    cheap    expensive
cheap:   3011     861
expensive:       1082     2160
kappa: 0.446
ConfusionMatrix:
True:    cheap    expensive
cheap:   3011     861
expensive:       1082     2160
AUC: 0.775 (positive class: expensive)
```

accuracy: 72.69%

	true cheap	true expensive	class precision
pred. cheap	3011	861	77.76%
pred. expensive	1082	2160	66.63%
class recall	73.56%	71.50%	

Check the model performance. What is its accuracy and kappa, can accuracy be trusted? (more on this later)

DEAKIN BUSINESS SCHOOL